



Deploying ECS autoscaling app with
Jenkins and Terraform



Konstantin Shcherban

Senior DevOps Engineer at  | GROUP

8 years of experience

Infrastructure, monitoring, deployments, etc.

About

Headquartered in Berlin, AUTO1 Group is Europe's leading car trading platform.


- Diverse team of over 3,500 employees
- More than 55 nationalities
- Over 1,000,000 private customers
- Over 45,000 professional partners.

The brand independent automotive company owns business units like AUTO1.com, wirkaufendeinauto.de and Autohero.




*The Co-Founders
Christian Bertermann
and Hakan Koç*

Managing service is tough




Managing microservices is tougher



What is AWS ECS?


Why should you use it?

AWS ECS




- black box container scheduler
- best integration with AWS services (even better than EKS)
- both managed infra *EC2* and serverless *Fargate*

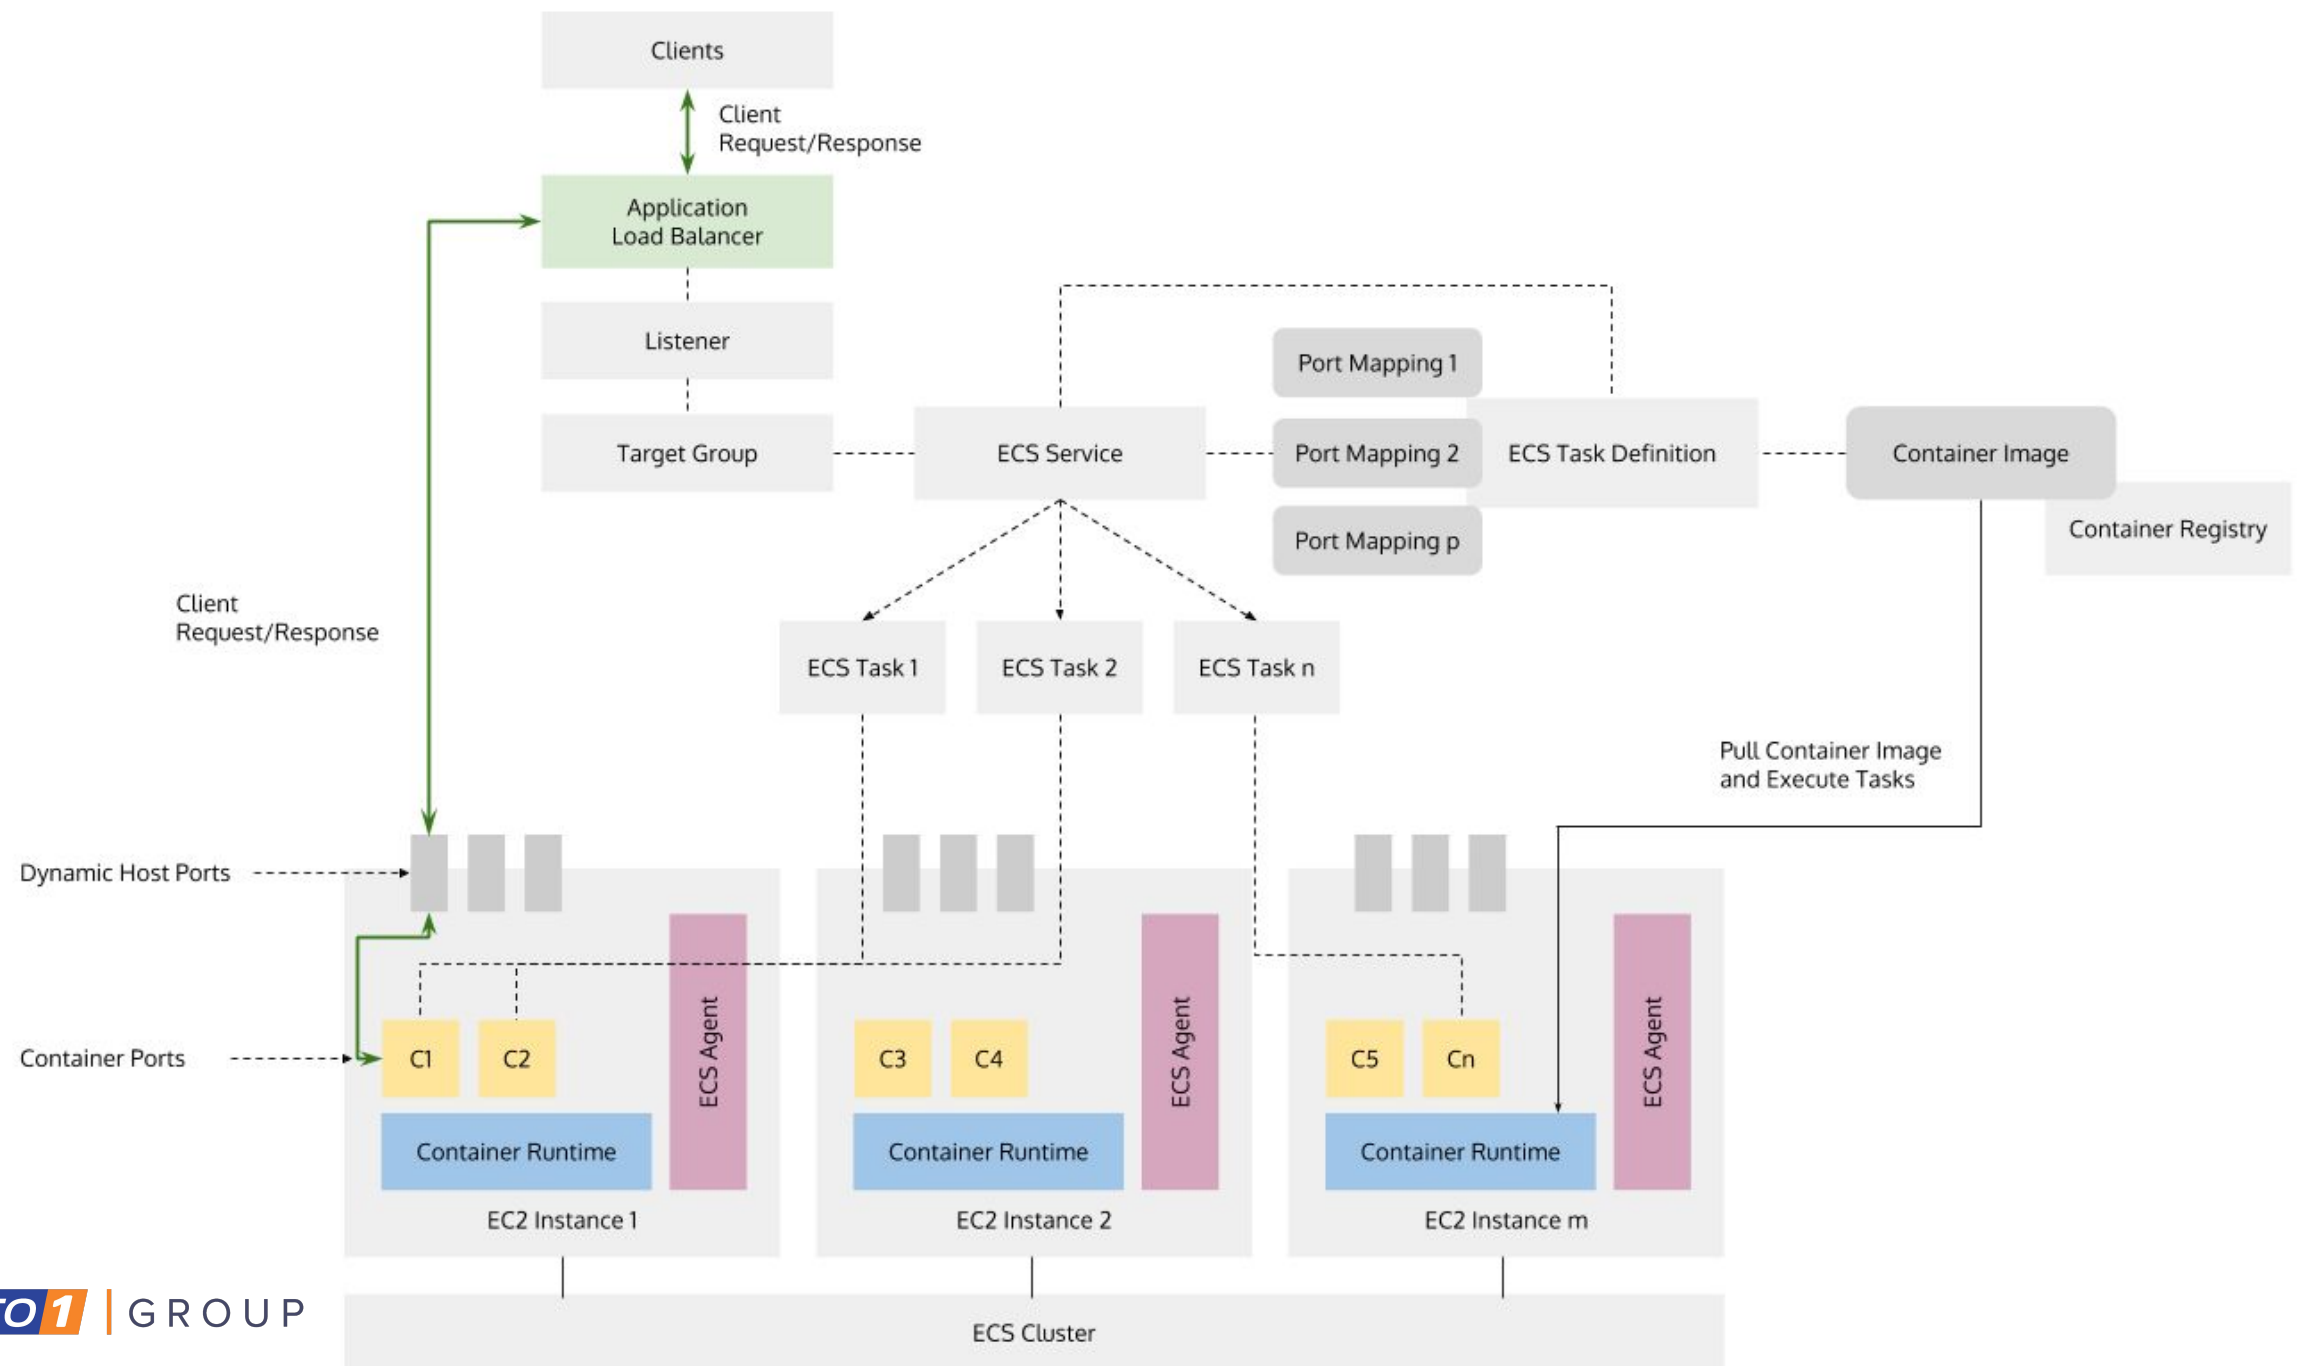
AWS ECS



- pay only for instances you manage, no weird pricing like EKS, scheduler is free
- **stable**, AWS uses it internally
- not many features unlike vanilla Kubernetes

ECS architecture





ECS services

- Daemon
- Replica
- Batch task


Do you really know AWS?

<https://goo.gl/qVkWNp>




Why to use ECS?

- ALB/NLB balancers
- route53 service discovery
- IAM task roles
- EBS/EFS volumes



Why to use ECS?

- Cloudwatch metrics/logs
- VPC networking, security groups per service
- ECS autoscaling
- ECR
- Serverless




Why *NOT* to use ECS?

- full control over cluster wanted
- features needed that ECS doesn't support
- fatal flaw — you didn't write it
(in case you have own container scheduler)

Autoscaling

- EC2 autoscaling
- ECS autoscaling



... they are **not** friends

Autoscaling

<https://aws.amazon.com/blogs/compute/how-to-automate-container-instance-draining-in-amazon-ecs/>



AWS Lambda



Deployment, Sir



+



HashiCorp
Terraform

What to do before deployment?

- think about service discovery
- configure logs and metrics collection
- determine where to store secrets



Service Discovery

- native ECS route53 SRV
- ALB/NLB as entrypoint
- Consul
- custom solution



Logs

- Cloudwatch \$\$\$
- Elasticsearch, Logstash, Kibana
- third-party SaaS
- custom solution



Secrets management

- Hashicorp Vault
- S3 bucket
- AWS Systems Manager Parameter Store
- AWS Secrets Manager
- environment variables in task definitions
- custom solution



Deployment flow



1. Jenkins runs CI tests
2. Jenkins creates docker image and pushes it to ECR
3. Terraform deploys image to QA environment
4. QA and Devs test new QA deployment
5. Terraform deploys the same image to PROD

Docker image tagging

- tag latest image with latest tag
- tag all images with short git commit hash (c656597)
- if commit contains git tag, apply the tag (v1.0.1)
- tag all images deployed to PROD with prefix (prod_)

ECR image policies

- apply IAM policies to restrict image access
- apply lifecycle policies
(AWS limits 1000 images per repo)
- apply sharing policies with other accounts if needed

ECR image policies

```
"rules": [  
  {  
    "rulePriority": 1,  
    "description": "Keep prod tags",  
    "selection": {  
      "tagStatus": "tagged",  
      "tagPrefixList": ["prod", "v", "latest"],  
      "countType": "imageCountMoreThan",  
      "countNumber": 100  
    },  
    "action": {  
      "type": "expire"  
    }  
  },  
],
```

Challenges



- Unification of services, especially between different teams
- Logs/metrics collection per service/container
- Security
- Alerting and resource tagging

Challenges

- LB has limit of 100 rules per listener
- update of large EC2 autoscaling group
- make cluster resource reservation == utilization



Demo time

github.com/auto1-oss/aws-ecs-jenkins-terraform



Thank you. Questions?

Konstantin Shcherban

Email: insider@prolinux.org

Github: github.com/kshcherban

Blog: www.prolinux.org (mostly in Russian)